

Neural Networks for NLP - Sheet 2

24.11.2025

Deadline: 8.12.2025 - 13:15, on OLAT as **pdf/.py/.ipynb** and/or after the lecture

Task 1: Analysing Pre-trained Word Embeddings

We will use pre-trained GloVe embeddings for analyzing contextualization of words. The pre-trained GloVe 6B embeddings can be downloaded [here](#).

- (a) Create two functions to compute the cosine and Euclidean distance between two vectors respectively. Use only Python's built-in functions and the math module.
- (b) Create a function that computes the most similar K words for a given input word. What are the most similar words to "awesome"?
- (c) Obtain the GloVe embeddings for the following words: red, orange, yellow, green, blue, purple, pink, brown, black, grey, white, violet, january, february, march, april, may, june, july, august, september, october, november, december. Calculate the cosine and euclidean distance between each pair of word embeddings and state your results. Do the word embeddings agree with your assumption of relatedness between words?
- (d) Create a function that computes analogies of type *man is to king like woman is to what?* Create ten other examples and check if the model supports your assumptions.

Task 2: Computing Word Embeddings – Word2Vec

- (a) Create a function that pre-processes and tokenizes the contents of the file *tutorial2.txt*. Then, produce a training set that consists of a list of center words and its surrounding context words.
- (b) Train a Neural Network to predict the center word given some context words (CBOW). Use a context window of size 1.
- (c) Train a Neural Network to predict the context words given some center word (Skip-gram). Use a context window of size 1.
- (d) Choose the best Word2Vec model from the previous questions (Skip-gram or CBOW) and justify your choice. Then, train the exact same model for context windows of 2 and 3. Does the context window affect the performance of the model? Explain in not more than 50 words.

Task 3: Implementing RNN-LSTM Classifier

If you read the file *tutorial2.txt* carefully, you might observe that each line represents a movie review, where every odd line is a positive review and the even ones are negative reviews.

- (a) Create a movie review dataset having all the positive and negative reviews from *tutorial2.txt*. Then, split the dataset into train (80%) and test (20%) sets. To obtain word embeddings of train and test splits, use the GloVe embeddings (used in Task 1). **Hint:** Don't forget to apply padding to your input sequences so that all the input sequences are of the same size.

- (b) Train a simple RNN-LSTM classifier using the layers defined in the *tutorial2_task3.py*. Use the trained model to obtain the classification accuracy on the test set and present the results.

In the Materials folder, you will find a file called *tutorial2_task3.py*. Use it as your base code.

Task 4: Resume Matching using Sentence Embeddings and PCA

In this task, you will work on matching resumes with a job query by using sentence embeddings and Principal Component Analysis (PCA) for visualization.

- (a) Load the dataset from the file *resumes_train.csv*. The dataset contains columns `resume` (the resume text) and `role` (the associated role). Use `pandas` to load the data.
- (b) Use a pre-trained model *SentenceTransformer (all-MiniLM-L6-v2)* to encode the `resume` text into numerical embeddings. Print the shape of the resulting embeddings.
- (c) Apply Principal Component Analysis (PCA) to reduce the embedding dimensions to 2. Plot the resumes in a 2D scatter plot where each resume is represented as a point colored by its `role`.
- (d) Define a job query of your choice (e.g., *"Data Engineer with Apache Airflow experience"*). Encode the job query into a numerical embedding and compute the Euclidean distance between the job query and each resume embedding.
- (e) Sort the resumes based on the computed distances and display the top 10 most relevant roles. Print the resume text for the closest match.